


Energy efficient computing EAR on Snellius

Benjamin Czaja
HPC Advisor SURF
June 2023


Energy-aware focus



Energy Aware Runtime (EAR) package provides an energy management framework for super computers.



EAR



Energy Aware Runtime (EAR)

Energy Aware Runtime (EAR) package provides an energy management framework for super computers.



- **EAR Node Manager (EARD) “Daemon”**

- Energy metrics via the RAPL (Running Average Power Limit) function
- Global energy limits or just offer global cluster monitoring

- **EARL** is a library that is loaded next to the application

- Offers application metrics monitoring
- Can select the frequencies based on the application behavior on the fly.
- Integrated with SLURM on Snellius.
- Intercepts the MPI symbols through the PMPI interface to provide “traces” of MPI applications.

```
#!/bin/bash

#SBATCH --ntasks=256
#SBATCH --time=24:0:0
#SBATCH -p thin --exclusive

#SBATCH --ear=on
#SBATCH --ear-policy=monitoring

module load 2021
module load foss/2021a
module load CMake/3.20.1-GCCcore-10.3.0
module load VASP6/6.2.1-foss-2021a-CUDA-11.3.1

srun vasp_std
```

Energy Aware Runtime (EAR)

Energy Aware Runtime (EAR) package provides an energy management framework for super computers.



- **EAR Node Manager (EARD) “Daemon”**

- Energy metrics via the RAPL (Running Average Power Limit) function
- Global energy limits or just offer global cluster monitoring

- **EARL** is a library that is loaded next to the application

- Offers application metrics monitoring
- Can select the frequencies based on the application behavior on the fly.
- Integrated with SLURM on Snellius.
- Intercepts the MPI symbols through the PMPI interface to provide “traces” of MPI applications.

```
#!/bin/bash
```

```
#SBATCH --ntasks=256
```

```
#SBATCH --time=24:0:0
```

```
#SBATCH -p thin --exclusive
```

```
#SBATCH --ear=on
```

```
#SBATCH --ear-policy=monitoring
```

```
module load 2021
```

```
module load foss/2021a
```



```
module load CMake/3.20.1-GCCcore-10.3.0
```

```
module load VASP6/6.2.1-foss-2021a-CUDA-11.3.1
```


```
srun vasp_std
```

Palabos: Lattice-Boltzmann Solver


Strong Scaling Benchmark




J. Latt, O. Malaspinas, D. Kontaxakis et al. / Computers and Mathematics with Applications 81 (2021) 334–350



initial constellation




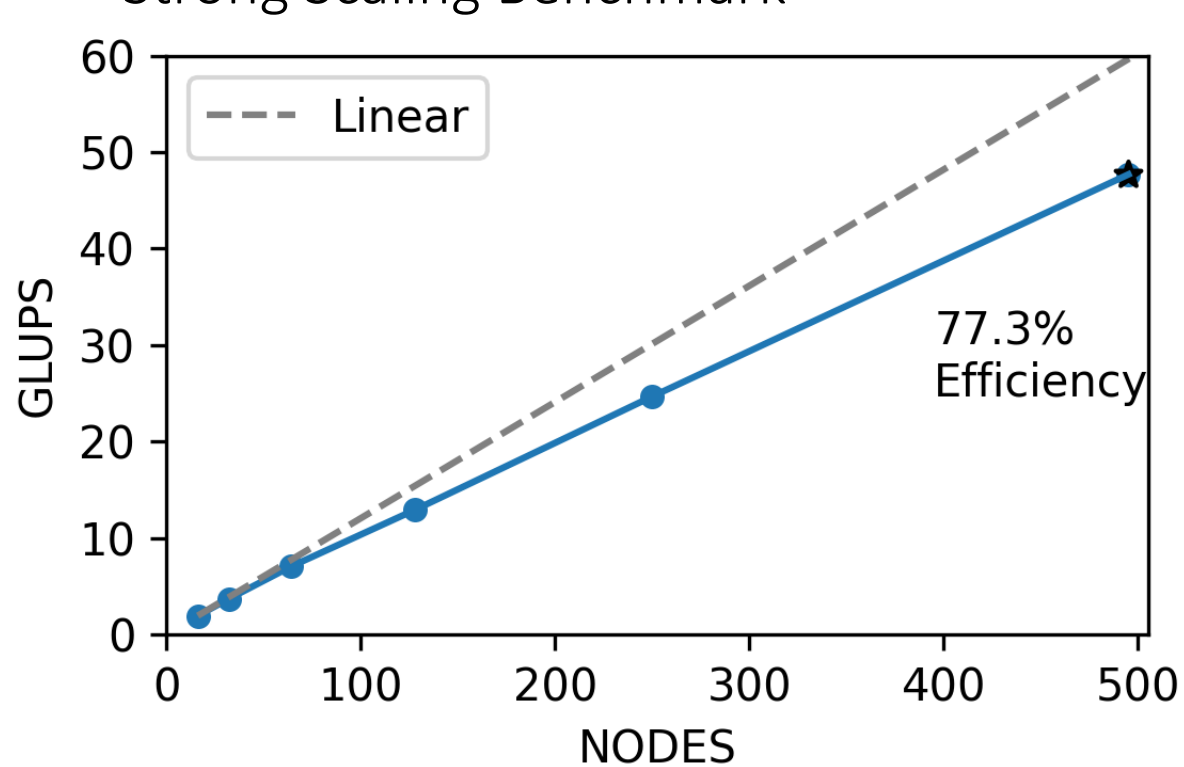
after collision kernel



after propagation, $t=t+1$

Palabos: Lattice-Boltzmann Solver

Strong Scaling Benchmark



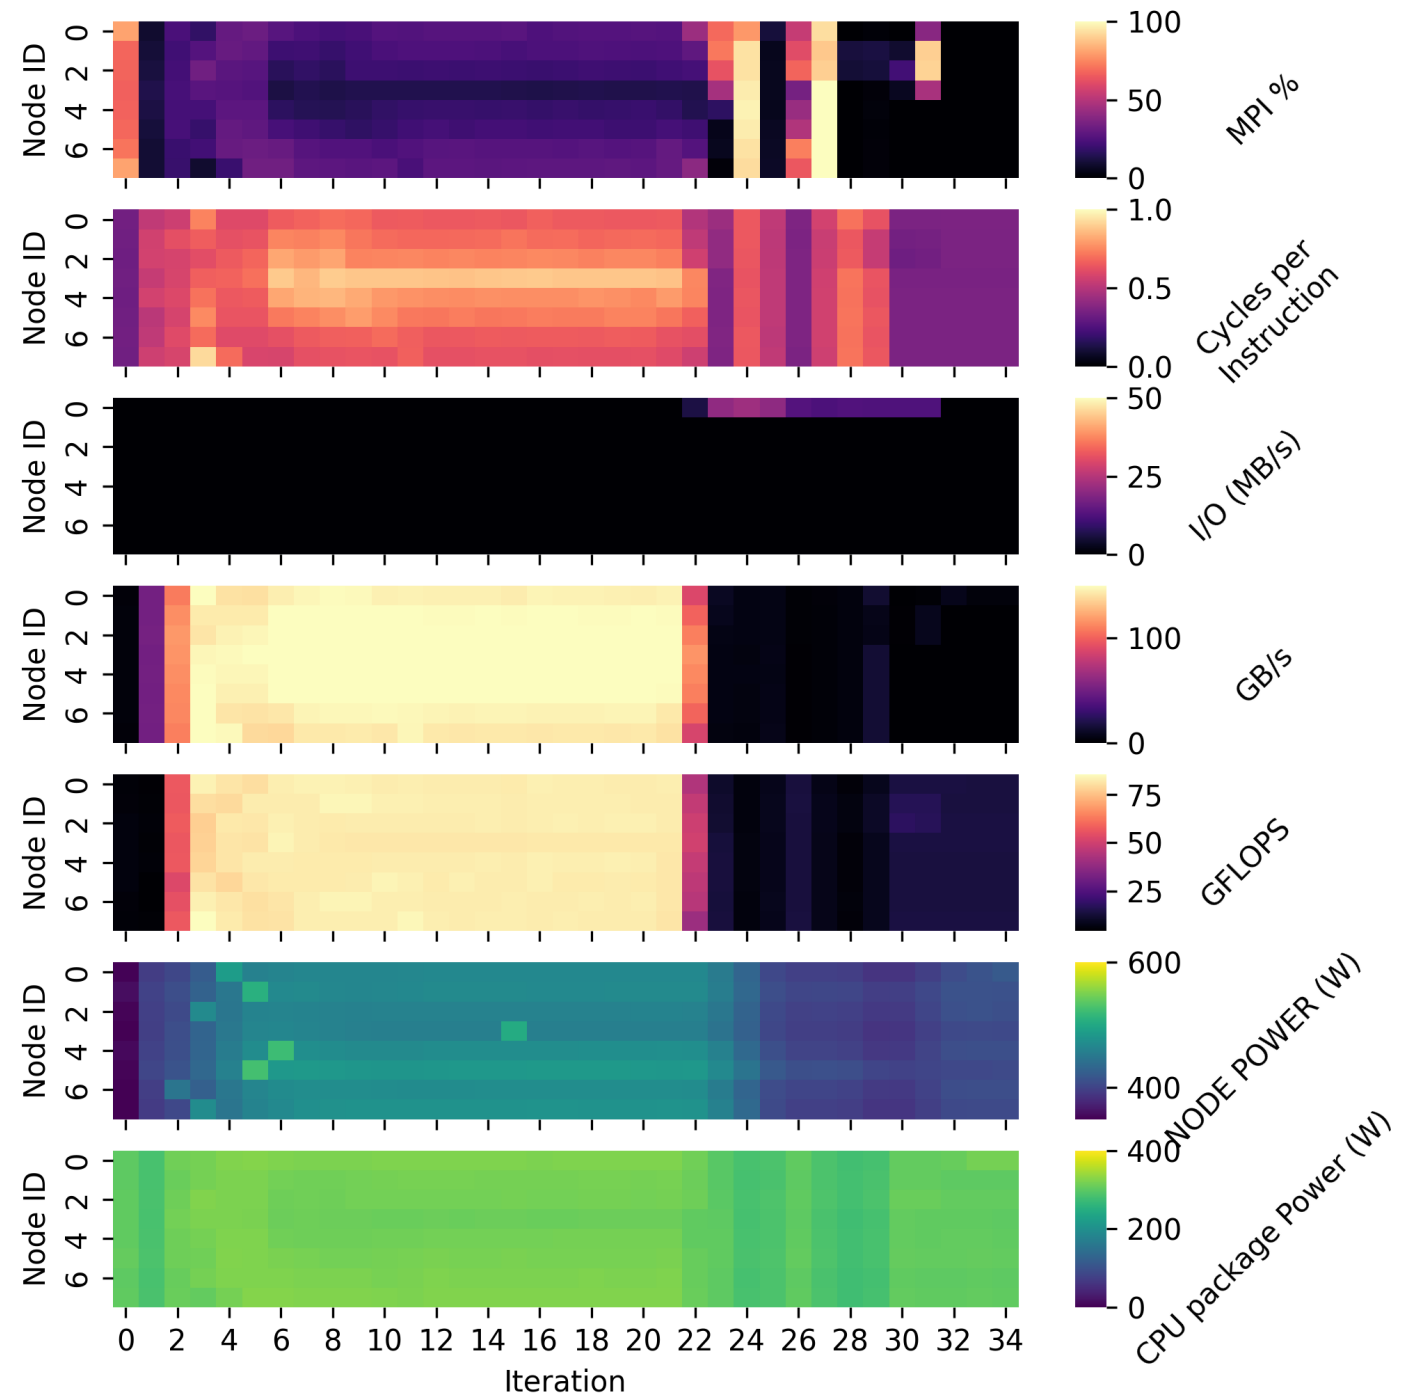
- An average electric car consumes about 0.2 kwh/km
- 495 node case (which ran for 9 minutes) just drove a car to Leeuwarden (149km from SP)!

Palabos:

Strong Scaling Benchmark

- Per-node, Per-iteration "traces"


- Node Power
- Avg CPU Freq (node)
- Main memory BW (GB/s)
- CPI (Cycles per instruction)
- MPI% (percentage spent in MPI calls)
- I/O (network communication)




Application Characterization

Surf Application Benchmark Suite (SABS)


- What aggregate metrics can we use to characterize an application?






Application characterization




Application characterization



Energy/Time to solution




Energy-aware focus

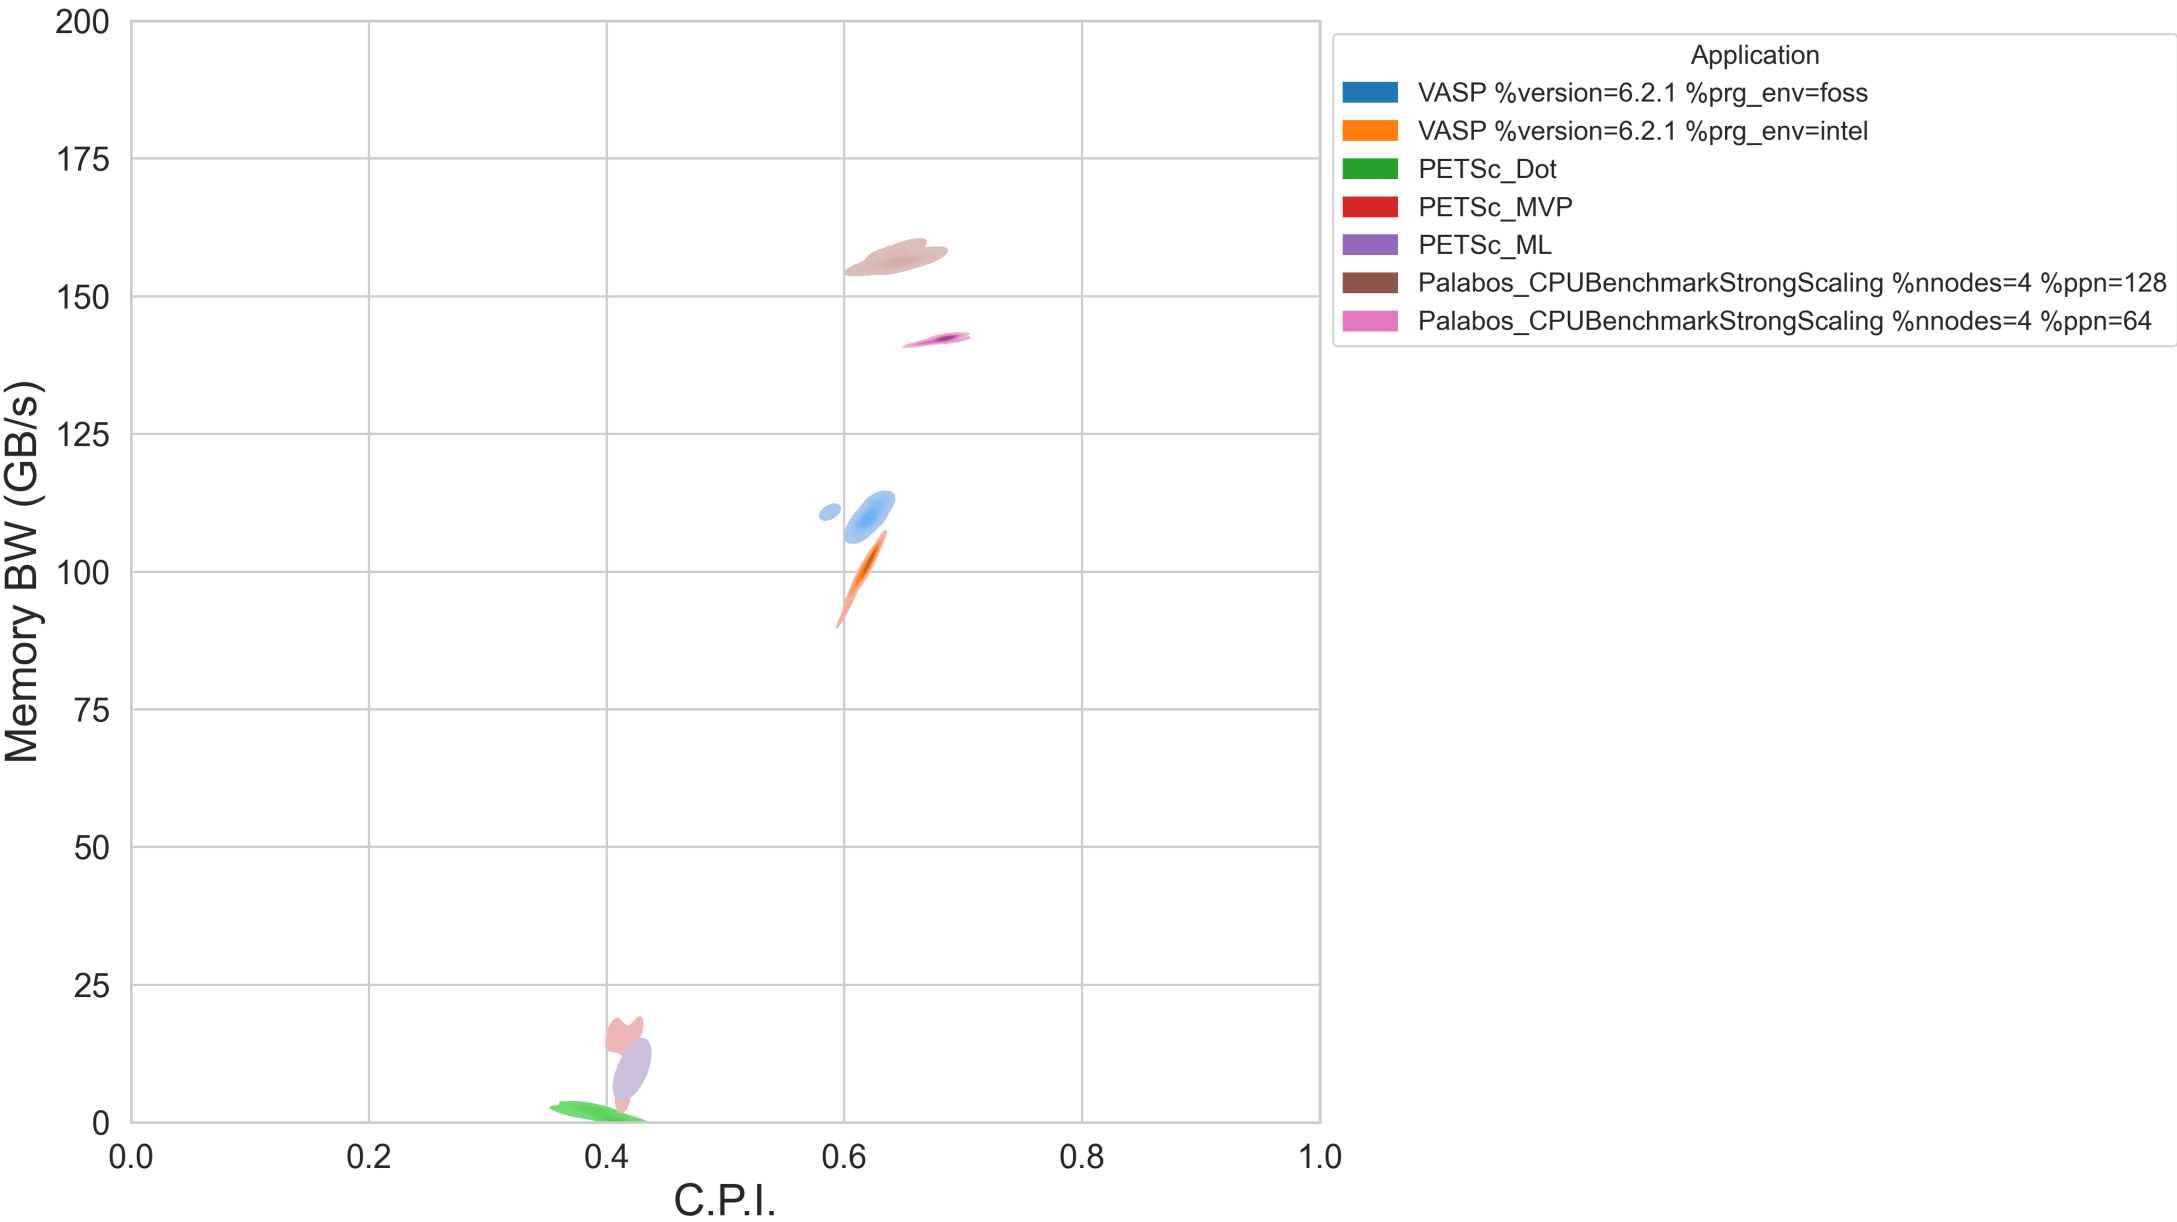


Application Optimization via EAR Policies

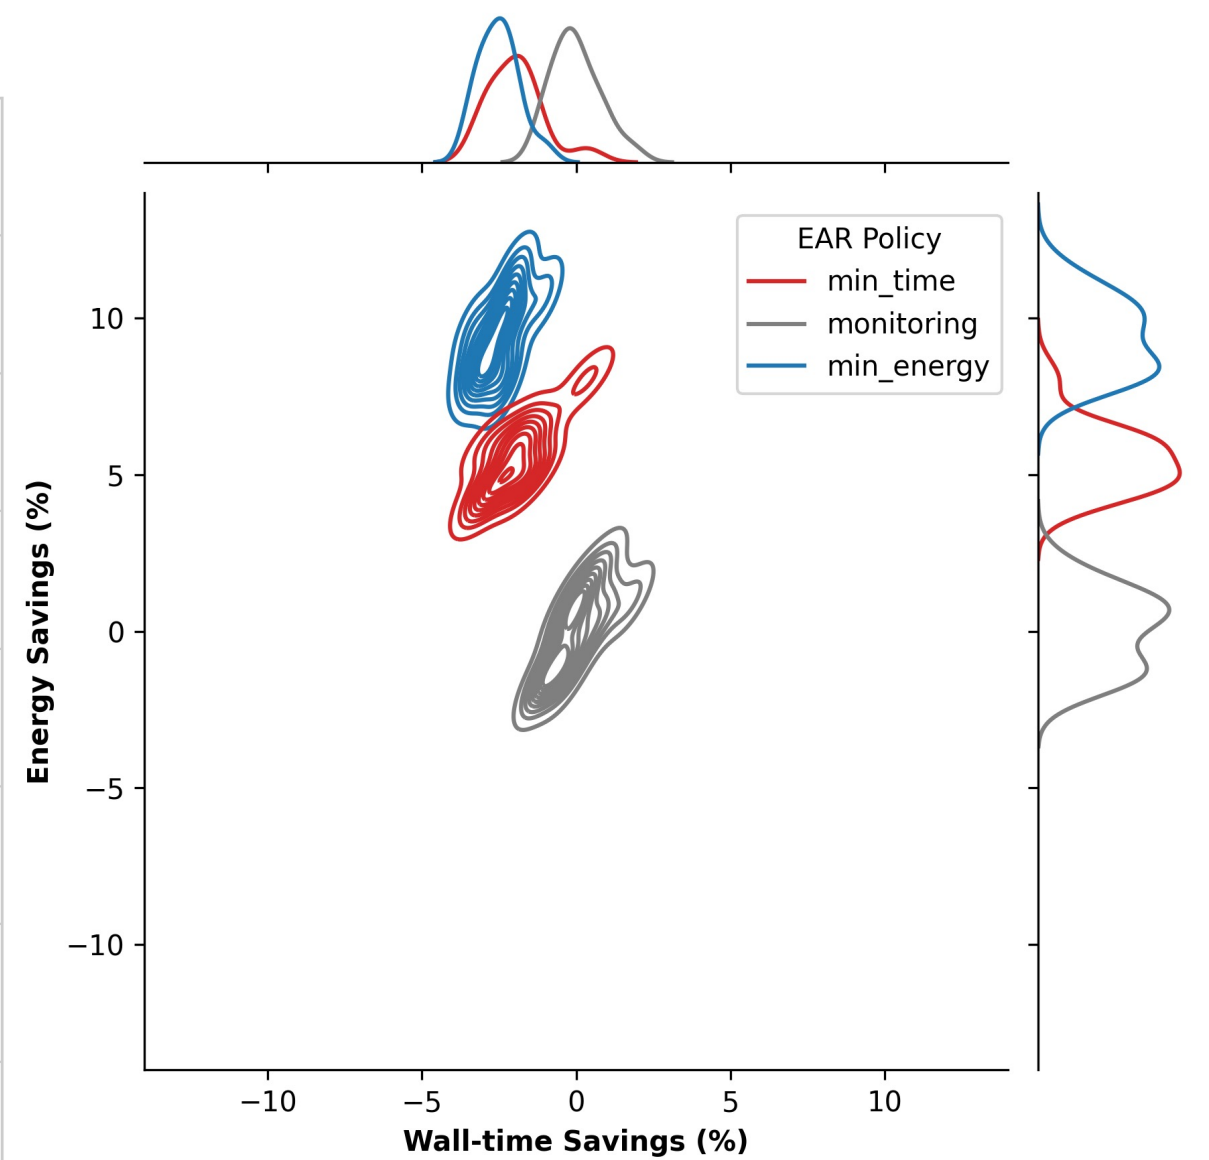
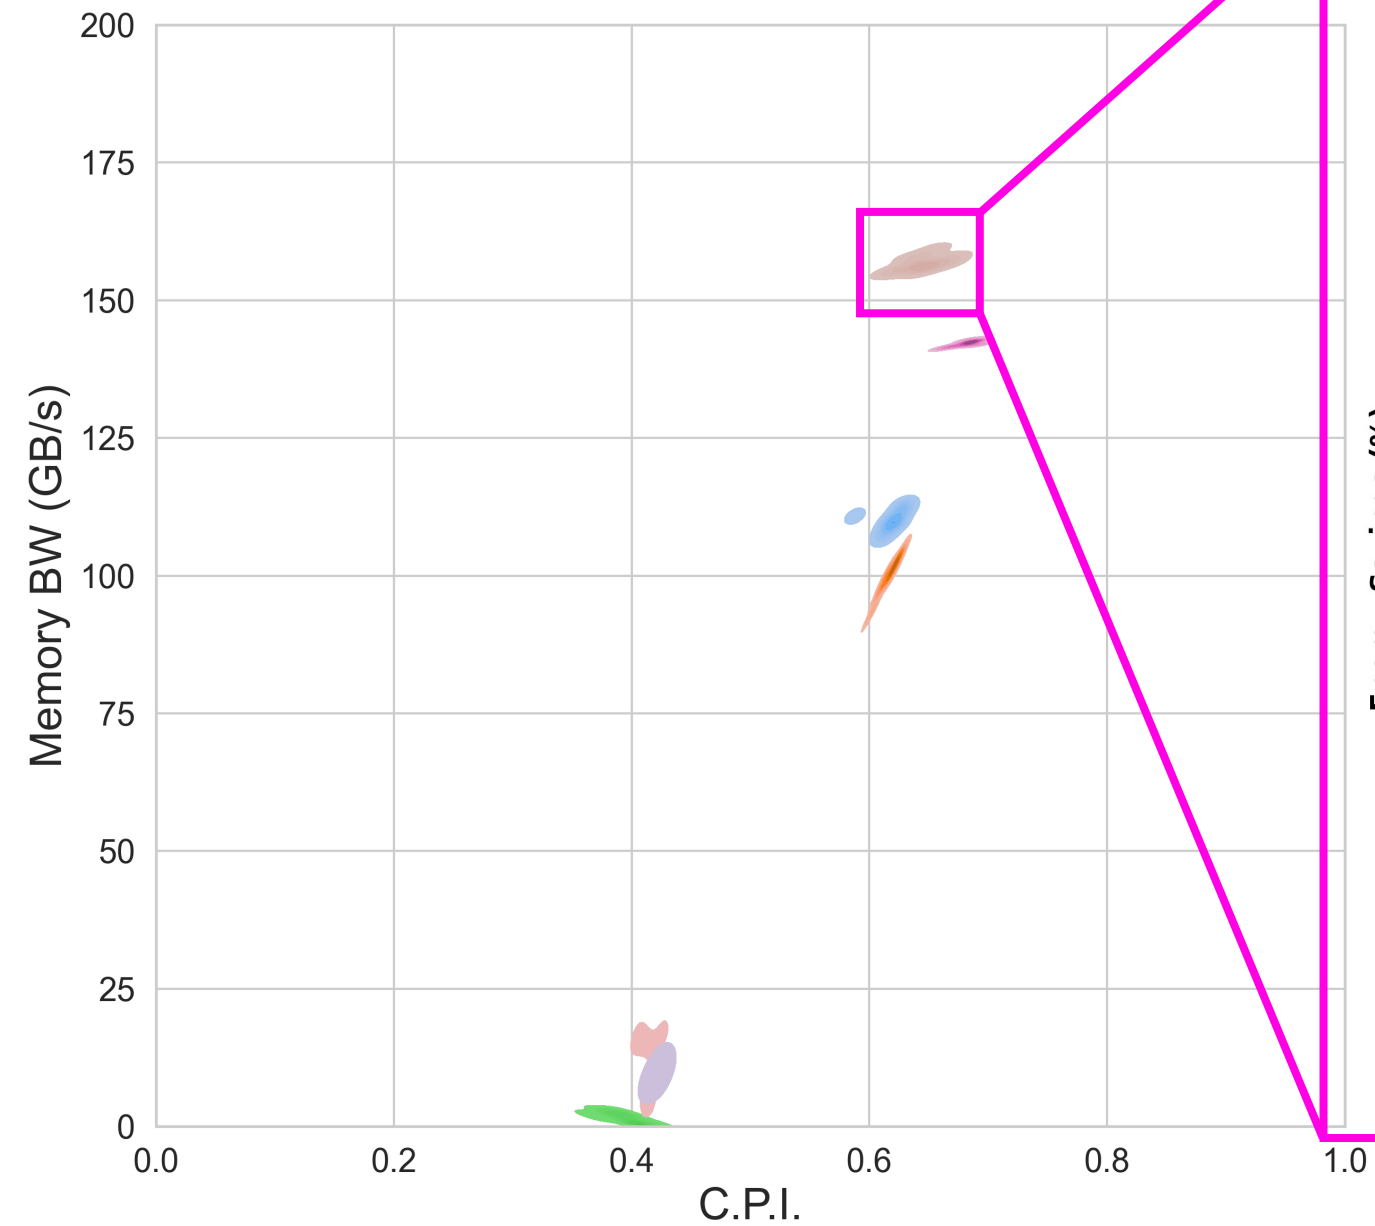
- monitoring
 - Does not affect the CPU Freq. Only monitors applications.
- min_energy
 - Memory intensive applications
 - Minimize the energy consumed with a limit to the performance degradation.
- min_time
 - CPU intensive applications
 - The goal of this policy is to improve the execution time while guaranteeing a minimum ratio between performance benefit and frequency increment that justifies the increased energy consumption from this frequency increment.



Application characterization



Application characterization/optimization



SURF - Energy Aware computing

- Contact us
 - <https://servicedesk.surf.nl/>
- EAR User Tutorial on Snellius
 - [https://servicedesk.surf.nl/wiki/display/WIKI/Energy+Aware+Runtime+\(EAR\)](https://servicedesk.surf.nl/wiki/display/WIKI/Energy+Aware+Runtime+(EAR))
- Energy Aware Runtime (EAR) --- Energy Aware Solutions
 - <https://www.eas4dc.com/solutions>
 - https://gitlab.bsc.es/ear_team/ear/-/wikis/home