

# Introduction to performance modeling

---

Xavier Álvarez-Farré

Training on energy-aware computing, June 16th, 2023.

High-Performance Computing Team, SURF, Science Park 140, 1098XG Amsterdam, The Netherlands



UNIVERSITY  
OF TWENTE.



東京工業大学  
Tokyo Institute of Technology

## Performance modeling

is the process of building analytical or empirical models to **predict and quantify the behavior** of applications on HPC systems.

- Aims to identify **performance bottlenecks** and optimize resource utilization.
- Enables understanding of the **impact of system configuration** and parameters on application performance.
- Facilitates **design decisions** for system upgrades or new architecture choices.

## Performance modeling

is the process of building analytical or empirical models to **predict and quantify the behavior** of applications on HPC systems.

- Aims to identify **performance bottlenecks** and optimize resource utilization.
- Enables understanding of the **impact of system configuration** and parameters on application performance.
- Facilitates **design decisions** for system upgrades or new architecture choices.

## Approaches for performance modeling

- **Analytical modeling** uses mathematical equations to represent the performance characteristics.
- **Empirical modeling** uses historical data or benchmarking to build models based on observed behavior.

## Performance modeling

is the process of building analytical or empirical models to **predict and quantify the behavior** of applications on HPC systems.

- Aims to identify **performance bottlenecks** and optimize resource utilization.
- Enables understanding of the **impact of system configuration** and parameters on application performance.
- Facilitates **design decisions** for system upgrades or new architecture choices.

## Approaches for performance modeling

- **Analytical modeling** uses mathematical equations to represent the performance characteristics.
- **Empirical modeling** uses historical data or benchmarking to build models based on observed behavior.

## Challenges in performance modeling

- Heterogeneous architectures and **complex interactions** between hardware and software components.
- **Scalability issues** when modeling large-scale systems with millions of concurrent processes.
- Accuracy and **validation of models** against empirical performance measurements.
- Model **portability** across different platforms and architectures.



Sports car



Family van



Transport truck



Sports car

- Power: 400 hp.
- Passengers: 2 pax.



Family van

- Power: 150 hp.
- Passengers: 7 pax.



Transport truck

- Power: 800 hp.
- Passengers: 3 pax.



Sports car

- Power: 400 hp.
- Passengers: 2 pax.
- Capacity: 0.5 m<sup>3</sup>.



Family van

- Power: 150 hp.
- Passengers: 7 pax.
- Capacity: 4 m<sup>3</sup>.



Transport truck

- Power: 800 hp.
- Passengers: 3 pax.
- Capacity: 45 m<sup>3</sup>.



Sports car

- Power: 400 hp.
- Passengers: 2 pax.
- Capacity:  $0.5 \text{ m}^3$ .
- Speed: 350 m/s.



Family van

- Power: 150 hp.
- Passengers: 7 pax.
- Capacity:  $4 \text{ m}^3$ .
- Speed: 210 m/s.



Transport truck

- Power: 800 hp.
- Passengers: 3 pax.
- Capacity:  $45 \text{ m}^3$ .
- Speed: 150 m/2.





Sports car

- Power: 400 hp.



Family van

- Power: 150 hp.



Transport truck

- Power: 800 hp.



Sports car

- Power: 400 hp.
- Passengers: 2 pax.
- Capacity:  $0.5 \text{ m}^3$ .
- Speed: 350 m/s.



Family van

- Power: 150 hp.
- Passengers: 7 pax.
- Capacity:  $4 \text{ m}^3$ .
- Speed: 210 m/s.



Transport truck

- Power: 800 hp.
- Passengers: 3 pax.
- Capacity:  $45 \text{ m}^3$ .
- Speed: 150 m/2.

## Application A

To deliver the maximum power possible.



Sports car

- Power: 400 hp.
- Passengers: 2 pax.
- Capacity:  $0.5 \text{ m}^3$ .
- Speed: 350 m/s.



Family van

- Power: 150 hp.
- Passengers: 7 pax.
- Capacity:  $4 \text{ m}^3$ .
- Speed: 210 m/s.



Transport truck

- Power: 800 hp.
- Passengers: 3 pax.
- Capacity:  $45 \text{ m}^3$ .
- Speed: 150 m/2.

## Application B

To win a speed race on a racing circuit.



Sports car

- Power: 400 hp.
- Passengers: 2 pax.
- Capacity:  $0.5 \text{ m}^3$ .
- Speed: 350 m/s.



Family van

- Power: 150 hp.
- Passengers: 7 pax.
- Capacity:  $4 \text{ m}^3$ .
- Speed: 210 m/s.




Transport truck

- Power: 800 hp.
- Passengers: 3 pax.
- Capacity:  $45 \text{ m}^3$ .
- Speed: 150 m/2.

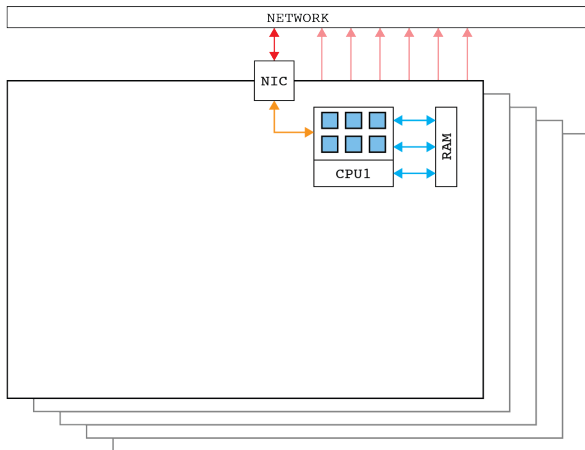
## Application C

To travel from point A to B with 5 passengers.

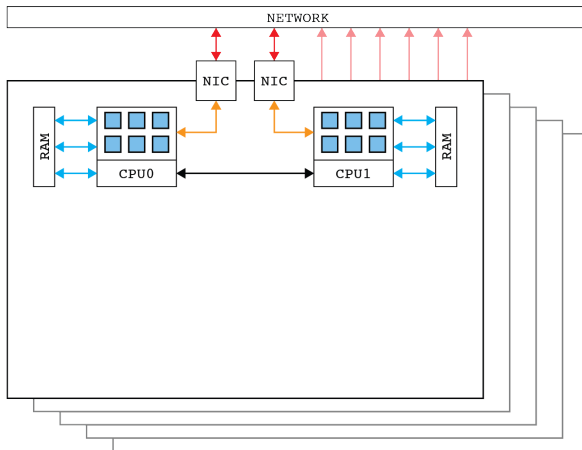
Modern HPC systems consist of **multiple hybrid computing nodes** interconnected via a communication infrastructure. The nodes are composed of **many hardware devices** of different architectures, such as central processing unit (**CPU**) or graphics processing unit (**GPU**), among others.



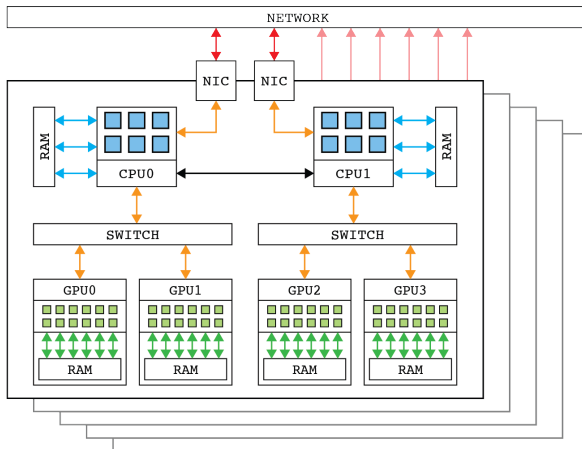
Modern HPC systems consist of **multiple hybrid computing nodes** interconnected via a communication infrastructure. The nodes are composed of **many hardware devices** of different architectures, such as central processing unit (**CPU**) or graphics processing unit (**GPU**), among others.



Modern HPC systems consist of **multiple hybrid computing nodes** interconnected via a communication infrastructure. The nodes are composed of **many hardware devices** of different architectures, such as central processing unit (**CPU**) or graphics processing unit (**GPU**), among others.




Modern HPC systems consist of **multiple hybrid computing nodes** interconnected via a communication infrastructure. The nodes are composed of **many hardware devices** of different architectures, such as central processing unit (**CPU**) or graphics processing unit (**GPU**), among others.









Consider computing kernels that perform a vast number of **independent operations** and **memory requests** which can be pipelined and parallelized. Distributed-memory parallelization requires large, **non-blocking point-to-point messages** between subsets of parallel processes that can **overlap** with bulk computations.



Consider computing kernels that perform a vast number of **independent operations** and **memory requests** which can be pipelined and parallelized. Distributed-memory parallelization requires large, **non-blocking point-to-point messages** between subsets of parallel processes that can **overlap** with bulk computations.





Consider computing kernels that perform a vast number of **independent operations** and **memory requests** which can be pipelined and parallelized. Distributed-memory parallelization requires large, **non-blocking point-to-point messages** between subsets of parallel processes that can **overlap** with bulk computations.



Parameter	Units	Description
$W$	flop	Number of operations
$Q$	byte	Main memory traffic
$Q_x$	byte	Network traffic



Consider computing kernels that perform a vast number of **independent operations** and **memory requests** which can be pipelined and parallelized. Distributed-memory parallelization requires large, **non-blocking point-to-point messages** between subsets of parallel processes that can **overlap** with bulk computations.



Parameter	Units	Description
$W$	flop	Number of operations
$Q$	byte	Main memory traffic
$Q_x$	byte	Network traffic



Parameter	Units	Description
$\pi$	flop/s	Peak performance
$\beta$	byte/s	Main memory bandwidth
$\beta_x$	byte/s	Network bandwidth

Consider computing kernels that perform a vast number of **independent operations** and **memory requests** which can be pipelined and parallelized.



Parameter	Units	Description
$W$	flop	Number of operations
$Q$	byte	Main memory traffic

Consider computing kernels that perform a vast number of **independent operations** and **memory requests** which can be pipelined and parallelized.



Parameter	Units	Description
$W$	flop	Number of operations
$Q$	byte	Main memory traffic

Parameter	Units	Description
$\pi$	flop/s	Peak performance
$\beta$	byte/s	Main memory bandwidth

The **axpy** kernel performing the linear combination of two vectors reads:

$$\mathbf{y} \leftarrow \alpha \mathbf{x} + \mathbf{y},$$

The **axpy** kernel performing the linear combination of two vectors reads:

$$\mathbf{y} \leftarrow \alpha \mathbf{x} + \mathbf{y},$$

where  $\alpha$  is a constant, and  $\mathbf{x}$  and  $\mathbf{y}$  are two vectors of size  $n$ . In such an operation, the processor performs:

- Work ( $W$ ):  $n$  additions and  $n$  multiplications, that is,  $2n$  floating-point operations.
- Memory traffic ( $Q$ ):  $2n$  elements read and  $n$  written, that is,  $3n \cdot 8$  bytes in double precision.




The **axpy** kernel performing the linear combination of two vectors reads:

$$\mathbf{y} \leftarrow \alpha \mathbf{x} + \mathbf{y},$$

where  $\alpha$  is a constant, and  $\mathbf{x}$  and  $\mathbf{y}$  are two vectors of size  $n$ . In such an operation, the processor performs:

- Work ( $W$ ):  $n$  additions and  $n$  multiplications, that is,  $2n$  floating-point operations.
- Memory traffic ( $Q$ ):  $2n$  elements read and  $n$  written, that is,  $3n \cdot 8$  bytes in double precision.

The kernel is illustrated as follows.




The **axpy** kernel performing the linear combination of two vectors reads:

$$\mathbf{y} \leftarrow \alpha \mathbf{x} + \mathbf{y},$$

where  $\alpha$  is a constant, and  $\mathbf{x}$  and  $\mathbf{y}$  are two vectors of size  $n$ . In such an operation, the processor performs:

- Work ( $W$ ):  $n$  additions and  $n$  multiplications, that is,  $2n$  floating-point operations.
- Memory traffic ( $Q$ ):  $2n$  elements read and  $n$  written, that is,  $3n \cdot 8$  bytes in double precision.

The kernel is illustrated as follows.



The time required to perform the calculations is:

$$t^W = \frac{W}{\pi} = \frac{2n}{\pi}.$$


The **axpy** kernel performing the linear combination of two vectors reads:

$$\mathbf{y} \leftarrow \alpha \mathbf{x} + \mathbf{y},$$

where  $\alpha$  is a constant, and  $\mathbf{x}$  and  $\mathbf{y}$  are two vectors of size  $n$ . In such an operation, the processor performs:

- Work ( $W$ ):  $n$  additions and  $n$  multiplications, that is,  $2n$  floating-point operations.
- Memory traffic ( $Q$ ):  $2n$  elements read and  $n$  written, that is,  $3n \cdot 8$  bytes in double precision.

The kernel is illustrated as follows.



The time required to perform the calculations is:

$$t^W = \frac{W}{\pi} = \frac{2n}{\pi}.$$

The time required to read and write the data is:

$$t^Q = \frac{Q}{\beta} = \frac{3n \cdot 8}{\beta}.$$


The **axpy** kernel performing the linear combination of two vectors reads:

$$\mathbf{y} \leftarrow \alpha \mathbf{x} + \mathbf{y},$$

where  $\alpha$  is a constant, and  $\mathbf{x}$  and  $\mathbf{y}$  are two vectors of size  $n$ . In such an operation, the processor performs:

- Work ( $W$ ):  $n$  additions and  $n$  multiplications, that is,  $2n$  floating-point operations.
- Memory traffic ( $Q$ ):  $2n$  elements read and  $n$  written, that is,  $3n \cdot 8$  bytes in double precision.

The kernel is illustrated as follows.



The time required to perform the calculations is:

$$t^W = \frac{W}{\pi} = \frac{2n}{\pi}.$$

The time required to read and write the data is:

$$t^Q = \frac{Q}{\beta} = \frac{3n \cdot 8}{\beta}.$$

The time required to execute the kernel:

$$t = \max(t^W, t^Q) = \max\left(\frac{2n}{\pi}, \frac{3n \cdot 8}{\beta}\right).$$

The `gemm` kernel performing the general matrix multiplication reads:

$$C \leftarrow \alpha AB + \beta C,$$

The `gemm` kernel performing the general matrix multiplication reads:

$$\mathbf{C} \leftarrow \alpha \mathbf{AB} + \beta \mathbf{C},$$

where  $\alpha$  and  $\beta$  are constants, and  $\mathbf{A} \in \mathbb{R}^{m \times k}$ ,  $\mathbf{B} \in \mathbb{R}^{k \times n}$ , and  $\mathbf{C} \in \mathbb{R}^{m \times n}$  are matrices. In such an operation, the processor performs:

- Work ( $W$ ):  $mn(k + 3)$  additions and multiplications, that is, approximately,  $2mnk$  floating-point operations.
- Memory traffic ( $Q$ ):  $mk + kn + mn$  elements read and  $mn$  written, that is,  $(mk + kn + 2mn) \cdot 8$  bytes in double precision.


The `gemm` kernel performing the general matrix multiplication reads:

$$\mathbf{C} \leftarrow \alpha \mathbf{A} \mathbf{B} + \beta \mathbf{C},$$

where  $\alpha$  and  $\beta$  are constants, and  $\mathbf{A} \in \mathbb{R}^{m \times k}$ ,  $\mathbf{B} \in \mathbb{R}^{k \times n}$ , and  $\mathbf{C} \in \mathbb{R}^{m \times n}$  are matrices. In such an operation, the processor performs:

- Work ( $W$ ):  $mn(k + 3)$  additions and multiplications, that is, approximately,  $2mnk$  floating-point operations.
- Memory traffic ( $Q$ ):  $mk + kn + mn$  elements read and  $mn$  written, that is,  $(mk + kn + 2mn) \cdot 8$  bytes in double precision.

The kernel is illustrated as follows.




The **gemm** kernel performing the general matrix multiplication reads:

$$\mathbf{C} \leftarrow \alpha \mathbf{A} \mathbf{B} + \beta \mathbf{C},$$

where  $\alpha$  and  $\beta$  are constants, and  $\mathbf{A} \in \mathbb{R}^{m \times k}$ ,  $\mathbf{B} \in \mathbb{R}^{k \times n}$ , and  $\mathbf{C} \in \mathbb{R}^{m \times n}$  are matrices. In such an operation, the processor performs:

- Work ( $W$ ):  $mn(k + 3)$  additions and multiplications, that is, approximately,  $2mnk$  floating-point operations.
- Memory traffic ( $Q$ ):  $mk + kn + mn$  elements read and  $mn$  written, that is,  $(mk + kn + 2mn) \cdot 8$  bytes in double precision.

The kernel is illustrated as follows.



The time required to perform the calculations is:

$$t^W = \frac{W}{\pi} = \frac{2mnk}{\pi}.$$




The **gemm** kernel performing the general matrix multiplication reads:

$$\mathbf{C} \leftarrow \alpha \mathbf{A} \mathbf{B} + \beta \mathbf{C},$$

where  $\alpha$  and  $\beta$  are constants, and  $\mathbf{A} \in \mathbb{R}^{m \times k}$ ,  $\mathbf{B} \in \mathbb{R}^{k \times n}$ , and  $\mathbf{C} \in \mathbb{R}^{m \times n}$  are matrices. In such an operation, the processor performs:

- Work ( $W$ ):  $mn(k + 3)$  additions and multiplications, that is, approximately,  $2mnk$  floating-point operations.
- Memory traffic ( $Q$ ):  $mk + kn + mn$  elements read and  $mn$  written, that is,  $(mk + kn + 2mn) \cdot 8$  bytes in double precision.

The kernel is illustrated as follows.



The time required to perform the calculations is:

$$t^W = \frac{W}{\pi} = \frac{2mnk}{\pi}.$$

The time required to read and write the data is:

$$t^Q = \frac{Q}{\beta} = \frac{8(mk + kn + 2mn)}{\beta}.$$


The **gemm** kernel performing the general matrix multiplication reads:

$$\mathbf{C} \leftarrow \alpha \mathbf{AB} + \beta \mathbf{C},$$

where  $\alpha$  and  $\beta$  are constants, and  $\mathbf{A} \in \mathbb{R}^{m \times k}$ ,  $\mathbf{B} \in \mathbb{R}^{k \times n}$ , and  $\mathbf{C} \in \mathbb{R}^{m \times n}$  are matrices. In such an operation, the processor performs:

- Work ( $W$ ):  $mn(k + 3)$  additions and multiplications, that is, approximately,  $2mnk$  floating-point operations.
- Memory traffic ( $Q$ ):  $mk + kn + mn$  elements read and  $mn$  written, that is,  $(mk + kn + 2mn) \cdot 8$  bytes in double precision.

The kernel is illustrated as follows.



The time required to perform the calculations is:

$$t^W = \frac{W}{\pi} = \frac{2mnk}{\pi}.$$

The time required to read and write the data is:

$$t^Q = \frac{Q}{\beta} = \frac{8(mk + kn + 2mn)}{\beta}.$$

The time required to execute the kernel:

$$t = \max\left(\frac{2mnk}{\pi}, \frac{8(mk + kn + 2mn)}{\beta}\right).$$

The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$

The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$

Application performance is often considered more relevant than elapsed time itself. Let's divide the work  $W$  by  $t^{RL}$ :

$$R^{RL} = \frac{W}{t^{RL}} = \min \left( \pi, \frac{W}{Q} \beta \right) = \min(\pi, I\beta).$$

The model introduces operational intensity,  $I$ , as the ratio of floating-point operations, in flops, to the memory traffic, in bytes.

The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$

Application performance is often considered more relevant than elapsed time itself. Let's divide the work  $W$  by  $t^{RL}$ :

$$R^{RL} = \frac{W}{t^{RL}} = \min \left( \pi, \frac{W}{Q} \beta \right) = \min(\pi, I\beta).$$

The model introduces operational intensity,  $I$ , as the ratio of floating-point operations, in flops, to the memory traffic, in bytes.

## Operational intensity

Fundamental characteristic of the algorithm and determines its theoretically achievable performance.

# The roofline model

The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$

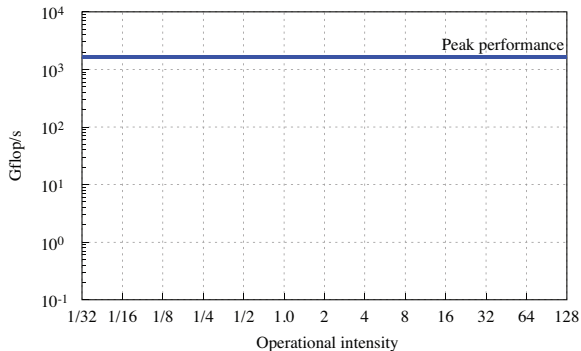
Application performance is often considered more relevant than elapsed time itself. Let's divide the work  $W$  by  $t^{RL}$ :

$$R^{RL} = \frac{W}{t^{RL}} = \min \left( \pi, \frac{W}{Q} \beta \right) = \min(\pi, I\beta).$$

The model introduces operational intensity,  $I$ , as the ratio of floating-point operations, in flops, to the memory traffic, in bytes.

## Operational intensity

Fundamental characteristic of the algorithm and determines its theoretically achievable performance.



The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$


Application performance is often considered more relevant than elapsed time itself. Let's divide the work  $W$  by  $t^{RL}$ :

$$R^{RL} = \frac{W}{t^{RL}} = \min \left( \pi, \frac{W}{Q} \beta \right) = \min(\pi, I\beta).$$

The model introduces operational intensity,  $I$ , as the ratio of floating-point operations, in flops, to the memory traffic, in bytes.

## Operational intensity

Fundamental characteristic of the algorithm and determines its theoretically achievable performance.



The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$


Application performance is often considered more relevant than elapsed time itself. Let's divide the work  $W$  by  $t^{RL}$ :

$$R^{RL} = \frac{W}{t^{RL}} = \min \left( \pi, \frac{W}{Q} \beta \right) = \min(\pi, I\beta).$$

The model introduces operational intensity,  $I$ , as the ratio of floating-point operations, in flops, to the memory traffic, in bytes.

## Operational intensity

Fundamental characteristic of the algorithm and determines its theoretically achievable performance.





The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$

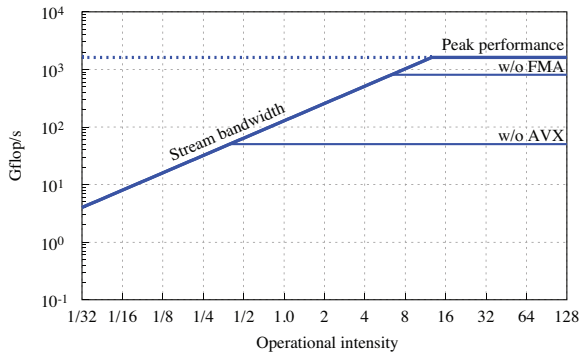
Application performance is often considered more relevant than elapsed time itself. Let's divide the work  $W$  by  $t^{RL}$ :

$$R^{RL} = \frac{W}{t^{RL}} = \min \left( \pi, \frac{W}{Q} \beta \right) = \min(\pi, I\beta).$$

The model introduces operational intensity,  $I$ , as the ratio of floating-point operations, in flops, to the memory traffic, in bytes.

## Operational intensity

Fundamental characteristic of the algorithm and determines its theoretically achievable performance.



The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$


Application performance is often considered more relevant than elapsed time itself. Let's divide the work  $W$  by  $t^{RL}$ :

$$R^{RL} = \frac{W}{t^{RL}} = \min \left( \pi, \frac{W}{Q} \beta \right) = \min(\pi, I\beta).$$

The model introduces operational intensity,  $I$ , as the ratio of floating-point operations, in flops, to the memory traffic, in bytes.

## Operational intensity

Fundamental characteristic of the algorithm and determines its theoretically achievable performance.



# The roofline model

The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$

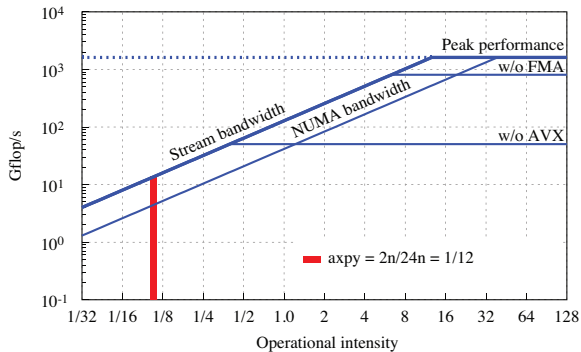
Application performance is often considered more relevant than elapsed time itself. Let's divide the work  $W$  by  $t^{RL}$ :

$$R^{RL} = \frac{W}{t^{RL}} = \min \left( \pi, \frac{W}{Q} \beta \right) = \min(\pi, I\beta).$$

The model introduces operational intensity,  $I$ , as the ratio of floating-point operations, in flops, to the memory traffic, in bytes.

## Operational intensity

Fundamental characteristic of the algorithm and determines its theoretically achievable performance.



# The roofline model

The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$


Application performance is often considered more relevant than elapsed time itself. Let's divide the work  $W$  by  $t^{RL}$ :

$$R^{RL} = \frac{W}{t^{RL}} = \min \left( \pi, \frac{W}{Q} \beta \right) = \min(\pi, I\beta).$$

The model introduces operational intensity,  $I$ , as the ratio of floating-point operations, in flops, to the memory traffic, in bytes.

## Operational intensity

Fundamental characteristic of the algorithm and determines its theoretically achievable performance.



# The roofline model

The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$


Application performance is often considered more relevant than elapsed time itself. Let's divide the work  $W$  by  $t^{RL}$ :

$$R^{RL} = \frac{W}{t^{RL}} = \min \left( \pi, \frac{W}{Q} \beta \right) = \min(\pi, I\beta).$$

The model introduces operational intensity,  $I$ , as the ratio of floating-point operations, in flops, to the memory traffic, in bytes.

## Operational intensity

Fundamental characteristic of the algorithm and determines its theoretically achievable performance.



# The roofline model

The elapsed time of a kernel running on a single computing unit is estimated as follows:

$$t^{RL} = \max \left( \frac{W}{\pi}, \frac{Q}{\beta} \right),$$


Application performance is often considered more relevant than elapsed time itself. Let's divide the work  $W$  by  $t^{RL}$ :

$$R^{RL} = \frac{W}{t^{RL}} = \min \left( \pi, \frac{W}{Q} \beta \right) = \min(\pi, I\beta).$$

The model introduces operational intensity,  $I$ , as the ratio of floating-point operations, in flops, to the memory traffic, in bytes.

## Operational intensity

Fundamental characteristic of the algorithm and determines its theoretically achievable performance.



Rank	Site	Computer	Cores	HPL Rmax (Pflop/s)	TOP500 Rank	HPCG (Pflop/s)	Fraction of Peak
1	RIKEN Center for Computational Science <b>Japan</b>	<b>Supercomputer Fugaku</b> — A64FX 48C 2.2GHz, Tofu interconnect D	7,630,848	442.01	2	16.00	3.0%
2	DOE/SC/Oak Ridge National Laboratory <b>United States</b>	<b>Frontier</b> — AMD Optimized 3rd Generation EPYC 64C 2GHz, Slingshot-11, AMD Instinct MI250X	8,730,112	1102.00	1	14.05	0.8%
3	EuroHPC/CSC <b>Finland</b>	<b>LUMI</b> — AMD Optimized 3rd Generation EPYC 64C 2GHz, Slingshot-11, AMD Instinct MI250X	2,220,288	309.10	3	3.408	0.8%
4	DOE/SC/Oak Ridge National Laboratory <b>United States</b>	<b>Summit</b> — IBM POWER9 22C 3.07GHz, Dual-rail Mellanox EDR Infiniband, NVIDIA Volta GV100	2,414,592	148.60	5	2.926	1.5%
5	EuroHPC/CINECA <b>Italy</b>	<b>Leonardo</b> — Xeon Platinum 8358 32C 2.6GHz, Quad-rail NVIDIA HDR100 Infiniband, NVIDIA A100 SXM4 64 GB	1,463,616	174.70	4	2.567	1.0%

## Performance modeling

is the process of building analytical or empirical models to **predict and quantify the behavior** of applications on HPC systems.

- Hardware architectures and system configurations are increasingly complex. Performance modeling aims to develop comprehensive and simplified machine models.



## Performance modeling

is the process of building analytical or empirical models to **predict and quantify the behavior** of applications on HPC systems.

- Hardware architectures and system configurations are increasingly complex. Performance modeling aims to develop comprehensive and simplified machine models.
- Energy and power constraints require understanding and optimizing application's performance.

## Performance modeling

is the process of building analytical or empirical models to **predict and quantify the behavior** of applications on HPC systems.

- Hardware architectures and system configurations are increasingly complex. Performance modeling aims to develop comprehensive and simplified machine models.
- Energy and power constraints require understanding and optimizing application's performance.
- Performance modeling provides insights into application behavior on HPC systems.

## Performance modeling

is the process of building analytical or empirical models to **predict and quantify the behavior** of applications on HPC systems.

- Hardware architectures and system configurations are increasingly complex. Performance modeling aims to develop comprehensive and simplified machine models.
- Energy and power constraints require understanding and optimizing application's performance.
- Performance modeling provides insights into application behavior on HPC systems.
- Understand your application's nature and requirements for accurate performance modeling.

## Performance modeling

is the process of building analytical or empirical models to **predict and quantify the behavior** of applications on HPC systems.

- Hardware architectures and system configurations are increasingly complex. Performance modeling aims to develop comprehensive and simplified machine models.
- Energy and power constraints require understanding and optimizing application's performance.
- Performance modeling provides insights into application behavior on HPC systems.
- Understand your application's nature and requirements for accurate performance modeling.
- Choose specific models for each application based on its characteristics.

## Performance modeling

is the process of building analytical or empirical models to **predict and quantify the behavior** of applications on HPC systems.

- Hardware architectures and system configurations are increasingly complex. Performance modeling aims to develop comprehensive and simplified machine models.
- Energy and power constraints require understanding and optimizing application's performance.
- Performance modeling provides insights into application behavior on HPC systems.
- Understand your application's nature and requirements for accurate performance modeling.
- Choose specific models for each application based on its characteristics.
- Analytical models, like the roofline model, identify performance bottlenecks and guide optimization efforts.

## Performance modeling

is the process of building analytical or empirical models to **predict and quantify the behavior** of applications on HPC systems.

- Hardware architectures and system configurations are increasingly complex. Performance modeling aims to develop comprehensive and simplified machine models.
- Energy and power constraints require understanding and optimizing application's performance.
- Performance modeling provides insights into application behavior on HPC systems.
- Understand your application's nature and requirements for accurate performance modeling.
- Choose specific models for each application based on its characteristics.
- Analytical models, like the roofline model, identify performance bottlenecks and guide optimization efforts.
- Identify parallelism, optimize resource utilization, and improve system efficiency through performance modeling.

Thank you for your attention